

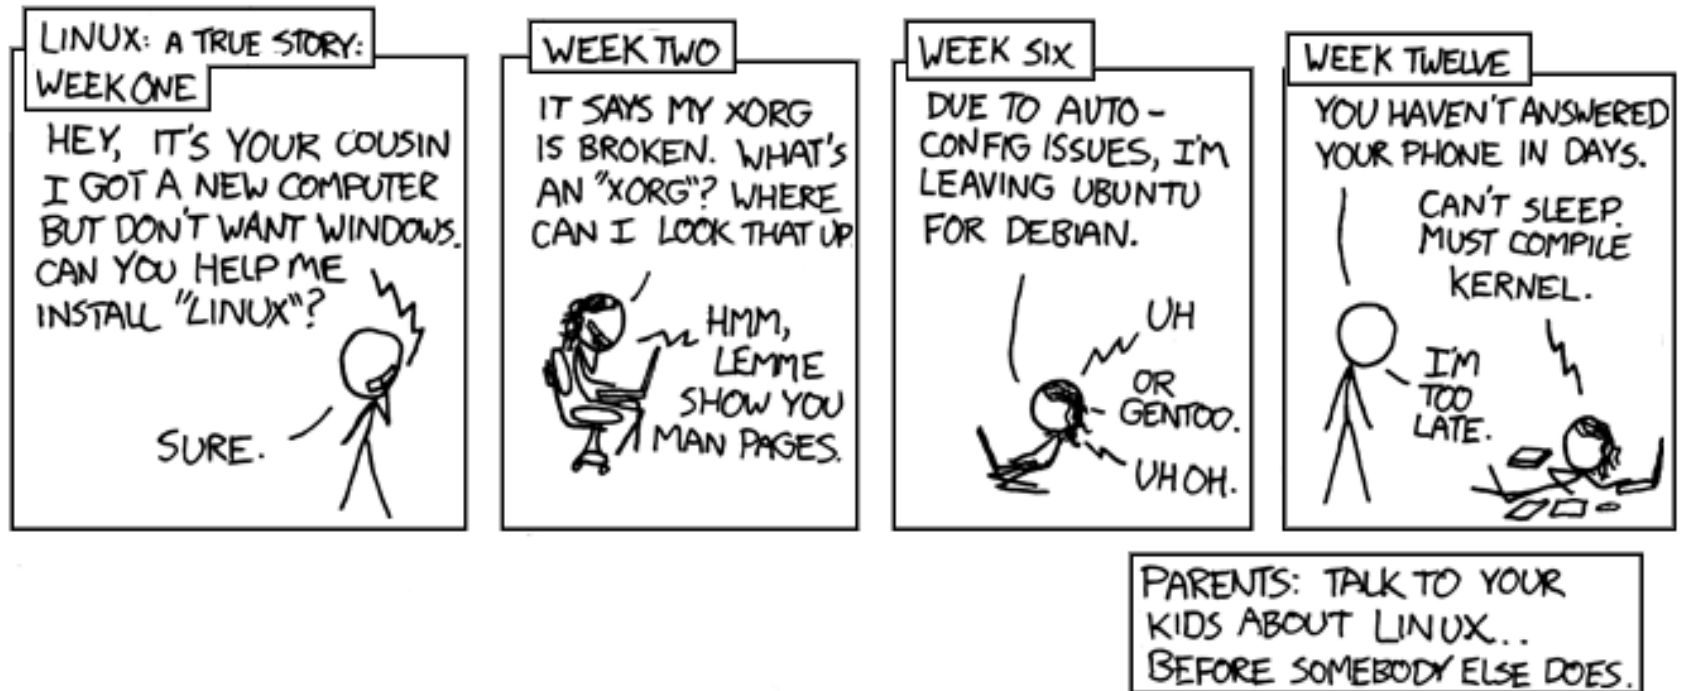
---

# Lecture 1

Introduction to Linux/Unix environment

slides created by Marty Stepp, modified by Jessica Miller (University of Washington)

# On to Linux



Courtesy XKCD.com

# Linux

---

- **Linux:** A kernel for a Unix-like operating system.
  - commonly seen/used today in servers, mobile/embedded devices, ...
- **GNU:** A "free software" implementation of many Unix-like tools
  - many GNU tools are distributed with the Linux kernel
- **distribution:** A pre-packaged set of Linux software.
  - examples: Ubuntu, Fedora
- key features of Linux:
  - **open source software:** source can be downloaded
  - free to use
  - constantly being improved/updated by the community



# Shell

---

- **shell**: An interactive program that uses user input to manage the execution of other programs.
  - `bash` : the default shell program on most Linux/Unix systems
- Why should I learn to use a shell when GUIs exist?

# Shell

---

- **shell**: An interactive program that uses user input to manage the execution of other programs.
  - `bash` : the default shell program on most Linux/Unix systems
- Why should I learn to use a shell when GUIs exist?
  - faster
  - work remotely
  - programmable
  - customizable
  - repeatable
- input, output, and errors
- directories: working/current directory, home directory

# Shell commands

---

command	description
exit	logs out of the shell
ls	lists files in a directory
pwd	outputs the current working directory
cd	changes the working directory
man	brings up the manual for a command

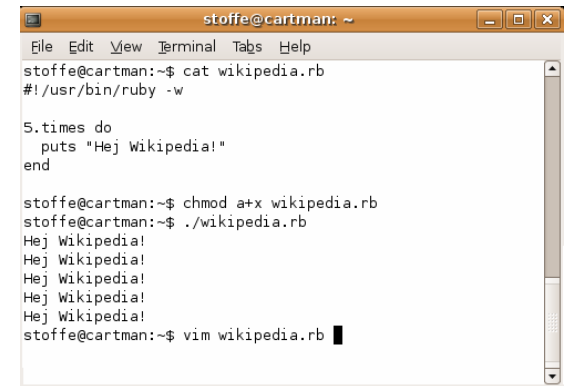
```
$ pwd
/homes/iws/dravir
$ cd CSE390
$ ls
file1.txt file2.txt
$ ls -l
-rw-r--r-- 1 dravir vgrad_cs 0 2010-03-29 17:45 file1.txt
-rw-r--r-- 1 dravir vgrad_cs 0 2010-03-29 17:45 file2.txt
$ cd ..
$ man ls
$ exit
```

# Relative directories

directory	description
.	the directory you are in ("working directory")
..	the parent of the working directory (../.. is grandparent, etc.)
~	your home directory (on many systems, this is /home/ <i>username</i> )
~ <i>username</i>	<i>username</i> 's home directory
~/Desktop	your desktop

# Shell commands

- many accept **arguments** or **parameters**
  - example: cp (copy) accepts a source and destination file path
- a program uses 3 streams of information:
  - stdin, stdout, stderr (standard in, out, error)
- **input**: comes from user's keyboard
- **output**: goes to console
- **errors** can also be printed (by default, sent to console like output)
- parameters vs. input
  - *parameters*: before Enter is pressed; sent in by shell
  - *input*: after Enter is pressed; sent in by user



```
stoffe@cartman: ~  
File Edit View Terminal Tabs Help  
stoffe@cartman:~$ cat wikipedia.rb  
#!/usr/bin/ruby -w  
  
5.times do  
  puts "Hej Wikipedia!"  
end  
  
stoffe@cartman:~$ chmod a+x wikipedia.rb  
stoffe@cartman:~$ ./wikipedia.rb  
Hej Wikipedia!  
Hej Wikipedia!  
Hej Wikipedia!  
Hej Wikipedia!  
Hej Wikipedia!  
Hej Wikipedia!  
stoffe@cartman:~$ vim wikipedia.rb
```



# Directory commands

---

command	description
<code>ls</code>	list files in a directory
<code>pwd</code>	output the current working directory
<code>cd</code>	change the working directory
<code>mkdir</code>	create a new directory
<code>rmdir</code>	delete a directory (must be empty)

- some commands (`cd`, `exit`) are part of the shell ("builtins")
- others (`ls`, `mkdir`) are separate programs the shell runs

# Command-line arguments

---

- most options are a - followed by a letter such as -c
  - some are longer words preceded by two - signs, such as --count
- parameters can be combined: `ls -l -a -r` can be `ls -lar`
- many programs accept a --help or -help parameter to give more information about that command (in addition to man pages)
  - or if you run the program with no arguments, it may print help info
- for many commands that accept a file name parameter, if you omit the parameter, it will read from standard input (your keyboard)
  - note that this can conflict with the previous tip

# Shell/system commands

---

command	description
man or info	get help on a command
clear	clears out the output from the console
exit	exits and logs out of the shell

command	description
date	output the system date
cal	output a text calendar
uname	print information about the current system

- "man pages" are a very important way to learn new commands  
man ls  
man man

# File commands

---

command	description
cp	copy a file
mv	move or rename a file
rm	delete a file
touch	create a new empty file, or update its last-modified time stamp

- caution: the above commands do not prompt for confirmation
  - easy to overwrite/delete a file; this setting can be overridden (how?)
- *Exercise* : Modify a `.java` file to make it seem as though you finished writing it on March 15 at 4:56am.

# Lecture summary

---

- Unix file system structure
- Commands for file manipulation, examination, searching
- Java compilation: using parameters, input, and streams
- Redirection and Pipes

# Unix file system

directory	description
/	root directory that contains all others (drives do not have letters in Unix)
/bin	programs
/dev	hardware devices
/etc	system configuration files <ul style="list-style-type: none"><li>▪ /etc/passwd stores user info</li><li>▪ /etc/shadow stores passwords</li></ul>
/home	users' home directories
/media, /mnt, ...	drives and removable disks that have been "mounted" for use on this computer
/proc	currently running processes (programs)
/tmp, /var	temporary files
/usr	user-installed programs

# File examination

---

command	description
<code>cat</code>	output a file's contents on the console
<code>more</code> or <code>less</code>	output a file's contents, one page at a time
<code>head</code> , <code>tail</code>	output the first or last few lines of a file
<code>wc</code>	count words, characters, and lines in a file
<code>du</code>	report disk space used by a file(s)
<code>diff</code>	compare two files and report differences

- Let's explore what we can do here...

# Searching and sorting

---

command	description
grep	search a file for a given string
sort	convert an input into a sorted output by lines
uniq	strip duplicate (adjacent) lines
find	search for files within a given directory
locate	search for files on the entire system
which	shows the complete path of a command

- **grep** is actually a very powerful search tool
- *Exercise* : Given a text file names.txt, display the students arranged by the reverse alphabetical order of their names.



# Keyboard shortcuts

---

**^KEY** means hold Ctrl and press **KEY**

key	description
Up arrow	repeat previous commands
Home/End or ^A/^E	move to start/end of current line
"	quotes surround multi-word arguments and arguments containing special characters
*	"wildcard" , matches any files; can be used as a prefix, suffix, or partial name
Tab	auto-completes a partially typed file/command name
^C or ^\	terminates the currently running process
^D	end of input; used when a program is reading input from your keyboard and you are finished typing
^Z	suspends (pauses) the currently running process
^S	don't use this; hides all output until ^Q is pressed

# Programming

---

command	description
<code>javac <i>ClassName</i>.java</code>	compile a Java program
<code>java <i>ClassName</i></code>	run a Java program
<code>python, perl, ruby, gcc, sml, ...</code>	compile or run programs in various other languages

- *Exercise* : Write/compile/run a program that prints "Hello, world!"

```
$ javac Hello.java
```

```
$ java Hello
```

```
Hello, world!
```

```
$
```

# Streams in the Shell

---

- Stdin, stdout, stderr
  - These default to the console
  - Some commands that expect an input stream will thus read from the console if you don't tell it otherwise.
- *Example:* `grep hi`
  - What happens? Why?

We can change the default streams to something other than the console via redirection.

# Output redirection

---

*command* > *filename*

- run *command* and write its output to *filename* instead of to console;
  - think of it like an arrow going from the command to the file...
  - if the file already exists, it will be overwritten (be careful)
  - >> appends rather than overwriting, if the file already exists
  - *command* > /dev/null suppresses the output of the command
- Example: `ls -l > myfiles.txt`
- Example: `java Foo >> Foo_output.txt`
- Example: `cat > somefile.txt`  
(writes console input to the file until you press ^D)

# Input redirection

---

*command* < *filename*

- run *command* and read its input from *filename* instead of console
  - whenever the program prompts the user to enter input (such as reading from a Scanner in Java), it will instead read the input from a file
  - some commands don't use this; they accept a file name as an argument
- Example: `java Guess < input.txt`
- Exercise: run hello world with the input stream as a file instead of the console
- Exercise: Also change the output stream to write the results to file
- again note that this affects *user input*, not *parameters*
- useful with commands that can process standard input or files:
  - e.g. `grep`, `more`, `head`, `tail`, `wc`, `sort`, `uniq`, `write`

# Combining commands

---

*command1* | *command2*

- run *command1* and send its console output as input to *command2*
- very similar to the following sequence:  
*command1* > *filename*  
*command2* < *filename*  
rm *filename*
- Examples: `diff students.txt names.txt | less`  
`sort names.txt | uniq`
- *Exercise* : `names.txt` contains CSE student first names, one per line. We are interested in students whose first names begin with "A", such as "Alisa".
  - Find out of how names beginning with "A" are in the file.
  - Then figure out how many characters long the name of the last student whose name starts with "A" is when looking at the names alphabetically.

# Commands in sequence

---

***command1 ; command2***

- run ***command1*** and then ***command2*** afterward (they are not linked)

***command1 && command2***

- run ***command1***, and if it succeeds, runs ***command2*** afterward
- will not run ***command2*** if any error occurs during the running of 1
- Example: Make directory songs and move my files into it.  
`mkdir songs && mv *.mp3 songs`

# Lecture summary

---

- A bit more on combining commands
- Processes and basic process management
- Text editors



# Review: Redirection and Pipes

---

- ***command > filename***
  - Write the output of ***command*** to ***filename*** (>> to append instead)
- ***command < filename***
  - Use ***filename*** as the input stream to ***command***
- ***command1 | command2***
  - Use the console output of ***command1*** as the input to ***command2***
- ***command1 ; command2***
  - Run ***command1*** and then run ***command2***
- ***command1 && command2***
  - Run ***command1***, if completed without errors then run ***command2***

# Tricky Examples

---

- The `wc` command can take multiple files: `wc names.txt student.txt`
  - Can we use the following to `wc` on every `txt` file in the directory?
    - `ls *.txt | wc`
- Amongst the top 250 movies in `movies.txt`, display the third to last movie that contains "The" in the title when movies titles are sorted.
- Find the disk space usage of the `man` program
  - Hints: use `which` and `du...`
  - Does `which man | du` work?

# The back-tick

---

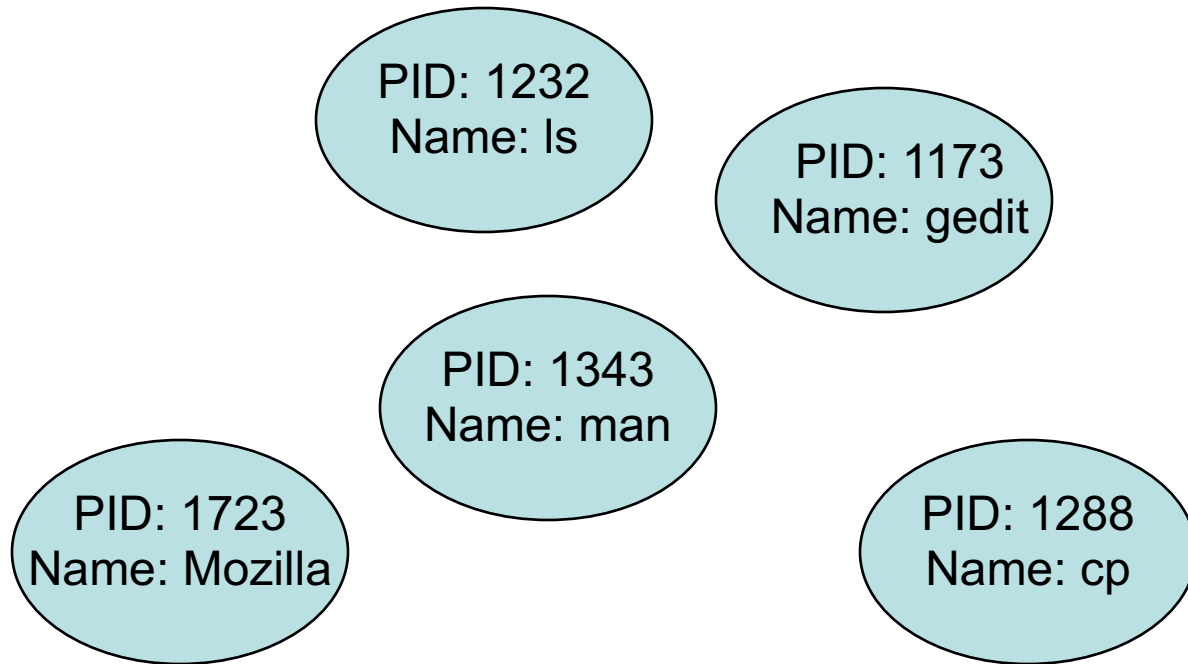
***command1*** ``command2``

- run ***command2*** and pass its console output to ***command1*** as a parameter; ``` is a back-tick, on the ~ key; not an apostrophe
- best used when ***command2***'s output is short (one line)
- Finish the example!
  - `du `which man``

# Processes

---

- **process:** a program that is running (essentially)
  - when you run commands in a shell, it launches processes for each
  - Process management is one of the major purposes of an OS



# Process commands

---

command	description
ps or jobs	list processes being run by a user; each process has a unique integer id (PID)
top	show which processes are using CPU/memory; also shows stats about the computer
kill	terminate a process by PID
killall	terminate several processes by name

- use `kill` or `killall` to stop a runaway process (infinite loop)
  - similar to `^C` hotkey, but doesn't require keyboard intervention

# Background processes

---

command	description
&	(special character) when placed at the end of a command, runs that command in the background
^Z	(hotkey) suspends the currently running process
fg, bg	resumes the currently suspended process in either the foreground or background

- If you run a graphical program like `gedit` from the shell, the shell will lock up waiting for the graphical program to finish
  - instead, run the program in the background, so the shell won't wait:  
`$ gedit resume.txt &`
  - if you forget to use `&`, suspend `gedit` with `^Z`, then run `bg`
  - lets play around with an infinite process...

# Connecting with ssh

---

command	description
ssh	open a shell on a remote server

- Linux/Unix are built to be used in multi-user environments where several users are logged in to the same machine at the same time
  - users can be logged in either locally or via the network
- You can connect to other Linux/Unix servers with ssh
  - once connected, you can run commands on the remote server
  - other users might also be connected; you can interact with them
  - can connect even from other operating systems

# Multi-user environments

---

command	description
whoami	outputs your username
passwd	changes your password
hostname	outputs this computer's name/address
w or finger	see info about people logged in to this server
write	send a message to another logged in user



# Network commands

---

command	description
<code>links</code> or <code>lynx</code>	text-only web browsers (really!)
<code>ssh</code>	connect to a remote server
<code>sftp</code> or <code>scp</code>	transfer files to/from a remote server (after starting sftp, use <code>get</code> and <code>put</code> commands)
<code>wget</code>	download from a URL to a file
<code>curl</code>	download from a URL and output to console
<code>pine</code> , <code>mail</code>	text-only email programs

# Text editors

---

command	description
<code>pico</code> or <code>nano</code>	simple but limited text editors
<code>emacs</code>	complicated text editor (powerful)
<code>vi</code> or <code>vim</code>	complicated text editor (powerful)

- most advanced Unix/Linux users learn `emacs` or `vi`
  - these editors are powerful but complicated and hard to learn
  - `nano` is simpler (hotkeys are shown on screen)

# Aliases

---

command	description
alias	assigns a pseudonym to a command

alias *name=command*

- must wrap the command in quotes if it contains spaces
- Example: When I type `q` , I want it to log me out of my shell.
- Example: When I type `ll` , I want it to list all files in long format.  

```
alias q=exit
```

```
alias ll="ls -la"
```
- *Exercise* : Make it so that typing `q` quits out of a shell.
- *Exercise* : Make it so that typing `woman` runs `man`.
- *Exercise* : Make it so that typing `seed` connects me to seed VM.

# Lecture summary

---

- Persistent settings for your bash shell
- User accounts and groups
- File permissions
- The Super User

# .bash\_profile and .bashrc

- Every time you log in to bash, the commands in `~/.bash_profile` are run
  - a `.` in front of a filename indicates a normally hidden file (`ls -a` to see)
  - you can put any common startup commands you want into this file
  - useful for setting up aliases and other settings for remote login
- Every time you launch a non-login bash terminal, the commands in `~/.bashrc` are run
  - useful for setting up persistent commands for local shell usage, or when launching multiple shells
  - often, `.bash_profile` is configured to also run `.bashrc`, but not always

# Users

---

*Unix/Linux is a multi-user operating system.*

- Every program/process is run by a user.
- Every file is owned by a user.
- Every user has a unique integer ID number (UID).
- Different users have different access permissions, allowing user to:
  - read or write a given file
  - browse the contents of a directory
  - execute a particular program
  - install new software on the system
  - change global system settings
  - ...

# Groups

---

command	description
groups	list the groups to which a user belongs
chgrp	change the group associated with a file

- **group:** A collection of users, used as a target of permissions.
  - a group can be given access to a file or resource
  - a user can belong to many groups
  - see who's in a group using `grep <groupname> /etc/group`
- Every file has an associated group.
  - the owner of a file can grant permissions to the group
- Every group has a unique integer ID number (GID).
- Exercise: create a file, see its default group, and change it

# File permissions

command	description
chmod	change permissions for a file
umask	set default permissions for new files

- *types* :      read (r),              write (w),              execute (x)
- *people* :      owner (u),              group (g),              others (o)

- on Windows, .exe files are executable programs;  
on Linux, any file with x permission can be executed
- permissions are shown when you type `ls -l`

*is it a directory?*

owner  
↓  
group  
↓  
others  
↓  
drwxrwxrwx



# Changing permissions

---

- letter codes: `chmod who(+-)what filename`

`chmod u+rw myfile.txt` (allow owner to read/write)

`chmod +x banner` (allow everyone to execute)

`chmod ug+rw,o-rwx grades.xls` (owner/group can read and

note: `-R` for recursive write; others nothing)

- octal (base-8) codes: `chmod NNN filename`

- three numbers between 0-7, for owner (u), group (g), and others (o)

- each gets +4 to allow read, +2 for write, and +1 for execute

`chmod 600 myfile.txt` (owner can read/write (rw))

`chmod 664 grades.dat` (owner rw; group rw; other r)

`chmod 751 banner` (owner rwx; group rx; other x)

# Exercises

---

- Change a file to grant full access (rwx) to everyone
  - Now change it to deny all access (rwx) from everyone
    - !!! is it dead?
    - I own this file. Can I change the user?

# Permissions don't travel

---

- Note in the previous examples that permissions are separate from the file
  - If I disable read access to a file, I can still look at its permissions
  - If I upload a file to a directory, its permissions will be the same as if I created a new file locally
- Takeaway: permissions, users, and groups reside on the particular machine you're working on. If you email a file or throw it on a thumbdrive, no permissions information is attached.
  - Why? Is this a gaping security hole?

# Lets combine things

---

- Say I have a directory structure, with lots of .txt files scattered
  - I want to remove all world permissions on all of the text files
  - First attempt:
    - `chmod -R o-rwx *.txt`
    - What happened?
  - Try and fix this using `find` and `xargs`!
    - `find -name "*.txt"`
    - `find -name "*.txt" | xargs chmod o-rwx`

# Super-user (root)

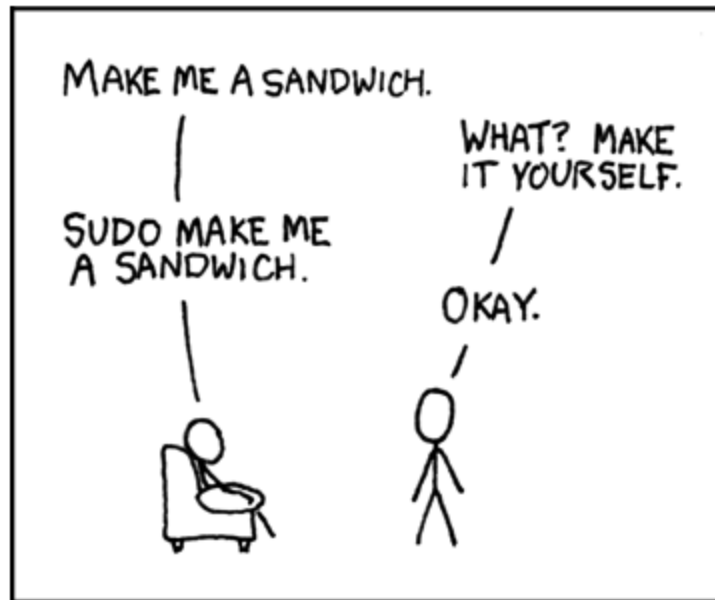
---

command	description
sudo	run a single command with root privileges (prompts for password)
su	start a shell with root privileges (so multiple commands can be run)

- **super-user:** An account used for system administration.
  - has full privileges on the system
  - usually represented as a user named root
- Most users have more limited permissions than root
  - protects system from viruses, rogue users, etc.
  - if on your own box, why ever run as a non-root user?
- Example: Install the `sun-java6-jdk` package on Ubuntu.  
`sudo apt-get install sun-java6-jdk`

# Playing around with power...

---



Courtesy XKCD.com

# Playing around with power...

---

- Create a file, remove all permissions
  - Now, login as root and change the owner and group to root
  - Bwahaha, is it a brick in a user's directory?
- Different distributions have different approaches
  - Compare Fedora to Ubuntu in regards to sudo and su...
- Power can have dangerous consequences
  - `rm *` might be just what you want to get rid of everything in a local directory
  - but what if you happened to be in `/bin...` and you were running as root...